

Take-Home Assessment

Overview

You're tasked with designing and implementing a data pipeline for a fictional e-commerce company, "ShopStream," that needs to analyze user behavior across multiple touchpoints to improve conversion rates and detect potential fraud.

Business Context

ShopStream receives data from multiple sources:

- Real-time clickstream events from web and mobile apps
- Order transactions from the payment system
- Product catalog updates from the inventory system
- Customer support interactions

The company needs to:

1. Create unified customer profiles
2. Calculate real-time metrics for fraud detection
3. Generate daily business intelligence reports
4. Enable data scientists to query historical data efficiently

The Task

Part 1: System Design (60% weight)

Design a data pipeline architecture that can handle the requirements above. Your design should include:

1. Architecture Diagram showing:
 - Data sources and ingestion methods
 - Processing layers (streaming/batch)
 - Storage solutions for different use cases
 - How different teams will access the data
2. Technology Choices with justification for:
 - Message queue/streaming platform
 - Processing frameworks
 - Storage systems (data lake, warehouse, etc.)
 - Orchestration tools
 - Data quality/monitoring solutions
3. Design Decisions addressing:
 - How you'll handle late-arriving data
 - Schema evolution strategy

- Data retention and archival policies
- Failure recovery mechanisms
- Scaling strategy for 10x growth

Part 2: Implementation (40% weight)

Implement a simplified version focusing on ONE of these scenarios (your choice):

Option A: Streaming Pipeline

- Process the provided sample clickstream data
- Implement sessionization (group events into user sessions with 30-minute inactivity timeout)
- Calculate real-time metrics: events per minute, unique users per minute
- Detect potential anomalies (e.g., >100 events from single user in 1 minute)

Option B: Batch Pipeline

- Process daily transaction and clickstream data
- Build user behavior aggregates (purchase patterns, browsing categories)
- Implement slowly changing dimension (SCD Type 2) for customer profiles
- Create data quality checks and quarantine bad records

Data Files Provided

You'll find the following sample data files in the data folder:

- `clickstream_events.json`: User interaction events from web/mobile
- `transactions.csv`: Completed and failed transactions
- `product_catalog.json`: Current product information
- `customer_support_interactions.json`: Support ticket data

Requirements & Constraints

Functional Requirements:

- Handle 1M events/minute during peak hours
- Sub-second latency for fraud detection
- 99.9% uptime for critical pipelines
- Support for GDPR compliance (data deletion requests)
- All timestamps are in UTC

Non-functional Requirements:

- Cost-conscious design (explain trade-offs)
- Monitoring and alerting strategy
- Data lineage tracking
- Development environment setup

Deliverables

1. System Design Document (PDF/Markdown):
 - Architecture diagram
 - Technology justifications
 - Design decisions with trade-offs
 - Capacity planning calculations
2. Code Implementation:
 - Working code for chosen scenario
 - README with setup instructions
 - Basic tests demonstrating functionality
 - Configuration for local development
3. Data Model:
 - Schema definitions for your chosen storage systems
 - Example queries for common use cases
4. Production Readiness (1-2 pages):
 - Deployment strategy
 - Monitoring/alerting plan
 - Disaster recovery approach
 - Performance optimization ideas

Time Expectation

This assessment should take 4-6 hours. We value quality over quantity - a well-thought-out partial solution is better than a rushed complete one.

Submission Instructions

1. Create a private GitHub repository
2. Commit your work with meaningful commit messages
3. Include all documentation in the repo
4. Share access with [interviewer email]

FAQ

Q: Should I implement everything in the design? A: No, implement only the chosen scenario. The design should be comprehensive, but the code should focus on demonstrating your implementation skills for one component.

Q: Can I use cloud-specific services? A: Yes, but explain why and provide cloud-agnostic alternatives.

Q: What if I'm unfamiliar with some technologies? A: Document your learning process and reasoning. We value problem-solving over memorization.

Q: Which programming languages can I use? A: Python is preferred, but Scala, Java, or Go are acceptable. Use what you're most comfortable with.

Q: Should I consider costs in my design? A: Yes! Include rough cost estimates and discuss trade-offs between cost and performance.